

# Off Policy Prediction

## Reinforcement Learning

- ◆ Target policy  $\pi$
- ◆ Behaviour policy  $\mu$
- ◆ Assumption of coverage: We require that wherever  $\pi(a | s) > 0$  implies  $\mu(a | s) > 0$ 
  - The behaviour policy  $\mu$  must be stochastic in states where it is not identical to  $\pi$
  - The target policy  $\pi$  may be deterministic

- ◆ Given a starting state  $S_t$ , the probability of the subsequent state-action trajectory,  $A_t, S_{t+1}, A_{t+1}, \dots, S_T$ , occurring under any policy  $\pi$  is

$$\sum_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)$$

◆ The importance sampling ratio is defined as:

$$\rho_t^T = \frac{\prod_{k=t}^{T-1} \pi(A_k|S_k) p(S_{k+1}|S_k, A_k)}{\prod_{k=t}^{T-1} \mu(A_k|S_k) p(S_{k+1}|S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{\mu(A_k|S_k)}$$

# Monte Carlo Algorithm

- ♦ The algorithm uses a batch of observed episodes following policy  $\mu$  to estimate  $v_\pi(s)$ .
- ♦ We define the set of all time steps in which state  $s$  is visited, denoted  $\mathcal{T}(s)$
- ♦ Let  $T(t)$  denote the first time of termination of episode following time  $t$ .
- ♦ Let  $G_t$  denote the return after  $t$  up through  $T(t)$ .
- ♦ Then  $\{G_t\}_{t \in \mathcal{T}(s)}$  are the returns that pertain to state  $s$ , and  $\{\rho_t^{T(t)}\}_{t \in \mathcal{T}(s)}$  are the corresponding sampling ratios

# Ordinary Importance Sampling

- ◆ To estimate  $v_\pi(s)$ , we simply scale the returns by the ratios and average the results

$$V(s) = \frac{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)} G_t}{|\mathcal{T}(s)|}.$$

- ◆ This is the averaging of scaled returns

# Weighted Importance Sampling


- ◆ This uses a weighted average of returns

$$V(s) = \frac{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)} G_t}{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)}},$$


- ◆ The simple average is not biased by can have high variance
- ◆ The weighted average has high bias, but has low variance.

- ♦ The variance of ordinary importance sampling is in general unbounded because the variance of the ratios is unbounded
- ♦ The variance of weighted importance-sampling estimator converges to zero even if the variance of the ratios themselves is infinite.

# Off-Policy Estimation of a Blackjack State Value




# Demonstration of Infinite Variance in Ordinary Importance Sampling



- ♦ The weighted importance-sampling algorithm would give an estimate of exactly 1 everafter the first episode that was consistent with the target policy (i.e., ended with the back action).
  - The algorithm produces a weighted average of the returns consistent with the target policy, all of which would be exactly 1.

# Ordinary Importance Sampling

- ◆ We can verify that the variance of the importance-sampling-scaled returns is infinite in this example.



$$\pi(\text{back}|s) = 1$$

$$\mu(\text{back}|s) = \frac{1}{2}$$

♦ Variance of a random variable  $X$

$$\text{Var}[X] = \mathbb{E}[(X - \bar{X})^2] = \mathbb{E}[X^2 - 2X\bar{X} + \bar{X}^2] = \mathbb{E}[X^2] - \bar{X}^2.$$

♦ The expectation of the square of the scaled returns:

$$\mathbb{E} \left[ \left( \prod_{t=0}^{T-1} \frac{\pi(A_t | S_t)}{\mu(A_t | S_t)} G_0 \right)^2 \right]$$

$$\mathbb{E}[X^2]$$

$$= \frac{1}{2} \cdot 0.1 \left( \frac{1}{0.5} \right)^2 \quad (\text{the length 1 episode})$$

$$+ \frac{1}{2} \cdot 0.9 \cdot \frac{1}{2} \cdot 0.1 \left( \frac{1}{0.5} \frac{1}{0.5} \right)^2 \quad (\text{the length 2 episode})$$

$$+ \frac{1}{2} \cdot 0.9 \cdot \frac{1}{2} \cdot 0.9 \cdot \frac{1}{2} \cdot 0.1 \left( \frac{1}{0.5} \frac{1}{0.5} \frac{1}{0.5} \right)^2 \quad (\text{the length 3 episode})$$

$$+ \dots$$

$$= 0.1 \sum_{k=0}^{\infty} 0.9^k \cdot 2^k \cdot 2$$

$$= 0.2 \sum_{k=0}^{\infty} 1.8^k$$

$$= \infty.$$

# Deep Reinforcement Learning

- ♦ How do we get feedback on our estimate  $Q(s, a)$  ?
- ♦ If we look one move ahead by taking action, we get an estimate of  $Q(s, a)$  as

$$R(s, a) + \gamma \max_{a'} Q(s', a')$$

- ◆ The squared difference between the old estimate and the new estimate can serve as the loss

$$\left[ Q(s, a) - \left( R(s, a) - \gamma \max_{a'} Q(s', a') \right) \right]^2$$

- ◆ This difference is also called as the temporal difference error or TD(0).
- ◆ If we looked two steps into the future it would be TD(1).


# Policy Gradient Methods

♦ Cart Pole

♦ State: (4 vector)

- Position of the cart
- Angle of the pole

(after the previous and  
the current move)



# Policy Gradient Methods

- ◆ The objective is to learn the policy, i.e., the recommended action for a given state.
- ◆ We play an entire game
  - We make 20 moves before the pole tips over
- ◆ We compute the discounted reward for the first state

$$D_0(\mathbf{s}, \mathbf{a}) = \sum_{t=0}^{n-1} \gamma^t R(s_t, a_t, s_{t+1})$$

# REINFORCE

- ♦ If we took  $n$  steps we compute the future discounted reward for any of the state-action combinations  $s_i, a_i$  from the recurrence relation

$$D_n(\mathbf{s}, \mathbf{a}) = 0$$

$$D_i(\mathbf{s}, \mathbf{a}) = R(s_i, a_i, s_{i+1}) + \gamma D_{i+1}(\mathbf{s}, \mathbf{a})$$

- ♦ Loss function

$$L(\mathbf{s}, \mathbf{a}) = \sum_{t=0}^{n-1} D_t(\mathbf{s}, \mathbf{a}) \left( -\log \Pr(a_t | s_t) \right)$$

# Actor-Critic Methods

- ◆ These work better than the actor programs
- ◆ They do not require playing the full episode

# Advantage Actor-Critic

- ◆ The advantage of a state-action pair is the difference between the state-action  $Q$  value and the state's value:

$$A(s, a) = Q(s, a) - V(s)$$

- ◆ Advantage is a negative number
- ◆ The a2c loss:

$$L_A(\mathbf{s}, \mathbf{a}) = \sum_{t=0}^{n-1} A(s_t, a_t) \left( \log \Pr(a_t | s_t) \right)$$

- ◆ The variance of advantage  $A(s_t, a_t)$  is smaller than the variance of  $D_t(\mathbf{s}, \mathbf{a})$
- ◆ It is easier to approximate a function with low variance than one with high.

# A2C

- ♦ How do we estimate  $A(s_t, a_t)$  ?
- ♦ We estimate  $Q(s, a)$  in the same way as we estimate  $D(\mathbf{s}, \mathbf{a})$
- ♦ We estimate  $V(s)$  by using another function called *Critic*.
  - It is another neural network trained to produce good estimates of  $V(s)$
  - It is trained by using the loss on the disparity between the actual rewards found and the output of the NN approximation.

# A2C

- ♦ Need not play an entire episode
- ♦ Can make use of multiple environments (games) as a batch of training examples.

# Experience Replay

- ♦ One big problem with the application of RL in real world is the acquisition of training data.
- ♦ In experience replay we use the same training data multiple times.
- ♦ For each time  $t$  we save  $\langle s_t, a_t, s_{t+1}, r_t \rangle$ .
  - We can do a forward and backward pass on our data.
  - We can play and replay, each time step in a random order.
  - Taking random actions from several different game plays reduces the correlation in data.

- ◆ But this too has a problem – the old data can go stale !
- ◆ We can instead keep a buffer
  - We then replace the oldest game in the buffer with a new game played using the new policy based upon the up-to-date parameters.